

Asset Pricing using Deep Neural Network

Team Members





Nikhil Katiki



Naveen Kumaran
Kumar

Agenda


- Problem Statement
- Why Neural network for this problem statement
- DNN parameters considered
- Base model
- Findings of Tuning parameters exploration (D1 60 models and D2 25 models)
- Best vs Worst Model comparison
- Out of box trials
- Learnings



Dataset 1



Problem Statement



Given

- A monthly stock price **returns for 690 months** (about 58 years) has been provided
- The dataset has **64 features** after pre-processing
- A total of **18719 stocks**

Goal


- To **predict the stock return accurately for 42 years** by creating the portfolio to sell bottom 10% stocks and buy top 10% stocks
- The model prediction is evaluated by Sharpe Ratio

Performance Evaluation (Sharpe Ratio)



- Sharpe ratio is a **metric for the investors to evaluate the return** of an investment considering its associated risk.
- **Sharpe Ratio = (Return of portfolio - Risk free rate)/Standard deviation of the portfolio's excess return**
- **Higher the sharpe ratio better** the investment in terms of return considering their associated risk

- Reference - <https://www.investopedia.com/terms/s/sharperatio.asp>







Why Neural network!!!



Deep Neural Network Parameters




DNN parameters considered



Sno.	Process	Rationale
1	Randomly	Since we wanted to take one direction
2	Layers	Increasing the number of layers increases prediction accuracy
3	Number of neurons	Increasing the number of neurons increases prediction accuracy
4	Activation functions	Functions to introduce the Non-linearity
5	Optimizer	Method to solve the objective function
6	Batch size	Updating weights using a batch
7	Learning Rate	Choosing an adaptive learning rate will help to “adapts” to the landscape
8	Epochs	Number of Epochs
9	Patience	Reduce the overfitting



Base Model – Sharpe ratio 5.12



Number of layers and neurons

Random tries from given base model								
Model No	Activation function	Batch size	Number of epochs	Optimizer	Number of layers	Network architecture	Patience	Sharpe Ratio
1	relu	128	50	adam (learning rate=0.001)	4	64,32,16,1	3	5.12
2	relu	128	80	adam (learning rate=0.001)	7	64,64,32,32,16,16,1	30	5.03
3	relu	128	80	adam (learning rate=0.001)	5	64,32,32,16,1	30	5.07

- Patience were increased significantly from 3 to allow the model to learn
- Learning – Comparing models with a change of multiple parameters is not an apple to apple comparison to make inference

No duplets with pyramid			
Model no.	Number of layers	Number of neurons	Sharpe Ratio
4	5	64,32,16,8,1	5.29
5	6	64,32,16,8,4,1	5.37
6	7	64,32,16,8,4,2,1	5.22

- The number of layers were increased considering the pyramid structure.
- Initially 64 were chosen as the cap of the structure since there are 64 features
- No duplets were considered initially

Number of layers and neurons


- Since it was a **regression problem** what if we add two neuron layer before the last layer!

1 Duplets with 2 neurons in last but one layer			
Model no.	Number of layers	Number of neurons	Sharpe Ratio
7	8	64,32,16,8,4,4,2,1	4.86
8	8	64,32,16,8,8,4,2,1	5.08
9	8	64,32,16,16,8,4,2,1	4.50
10	8	64,32,32,16,8,4,2,1	5.30
11	8	64,64,32,16,8,4,2,1	5.61

1 Duplets without 2 neurons in last but one layer			
Model no.	Number of layers	Number of neurons	Sharpe Ratio
12	7	64,32,16,8,4,4,1	4.68
13	7	64,32,16,8,8,4,1	5.29
14	7	64,32,16,16,8,4,1	5.44
15	7	64,32,32,16,8,4,1	5.18
16	7	64,64,32,16,8,4,1	5.25

- Learning** – Roughly there was an increasing trend of sharpe ratio when the duplet moved towards initial layers with higher neurons for both cases

Number of layers and neurons




2 Duplets without 2 neurons in last but one layer

Model no.	Number of layers	Number of neurons	Sharpe Ratio
17	8	64,64,32,32,16,8,4,1	5.12
18	8	64,64,32,16,16,8,4,1	4.99
19	8	64,64,32,16,8,8,4,1	5.29

3 Duplets without 2 neurons in last but one layer

Model no.	Number of layers	Number of neurons	Sharpe Ratio
20	9	64,64,32,32,16,16,8,4,1	5.15

- **Learning** – From all the **20 different model architectures**, we inferred that **adding more layers produced relatively better results**
- From running time, **we realized that it is not computationally feasible to add more layers** and tune other parameters.



Optimizer

Model no.	Optimizer	Sharpe Ratio
21	rmsprop(0.001)	5.22
22	SGD(learning_rate=0.001, momentum=0.0, nesterov=False)	1.26
23	SGD(learning_rate=0.001, momentum=0.0, nesterov=True)	0.77
24	SGD(learning_rate=0.001, momentum=0.9, nesterov=False)	1.92
25	SGD(learning_rate=0.001, momentum=0.9, nesterov=True)	2.43
26	Adagrad(learning_rate=0.001)	1.89
27	Adadelta(learning_rate=0.001)	0.98
28	Adam(learning_rate=0.001)	5.61

- **Learning** – When compared to SGD, momentum/Nesterov accelerated gradient, rms prop, Adagrad and Adadelta, Adam performed significantly better.
- **Rms prop** was the next high performing optimizer

Activation function

Model no.	Activation function	Sharpe Ratio
29	tanh	5.11
30	prelu	3.99
31	leaky relu	5.19
32	relu	5.61
33	relu, tanh	5.20
34	relu, prelu	5.18
35	swish	5.39
36	elu	4.97
37	sigmoid	5.16
38	relu last but one linear	5.07

Learnings

- **Basic activation functions** like **tanh** and **sigmoid** were tried
- Apart from **relu** from various references, other proven better performing activation functions like **leaky relu**, **prelu**, **elu**, **swish** were tried – **prelu** performed significantly **bad** and **swish** performed significantly **better**, a **common disadvantage of computationally expensive is notable**
- **Alternate activation functions** were tried with 1)**relu, one better performing as per literature (prelu)** 2) **relu, another average performing activation function (tanh)** - **No significant differences** were observed.
- **Relu except linear activation function in last but one layer** – Did not perform better
- Reference - <https://towardsdatascience.com/a-quick-guide-to-activation-functions-in-deep-learning-4042e7add5b>

Batch size

Model no.	Batchsize	Sharpe Ratio
39	64	4.93
45	80	5.17
43	100	4.85
47	128	5.61
44	150	4.39
46	200	5.18
40	256	4.58
41	512	4.78
42	1024	3.76

- Various batch size ranging from 64 to 1024 were tried

Learnings


- Increasing the batch size after a threshold decreased the sharpe ratio as expected.
- Sharpe ratio were fluctuating even with a batch size near to the high performing value (128). So might need multiple trials to find optimal batch size

Learning rate

Model no.	Learning rate	Sharpe Ratio
48	Adam(learning_rate=0.0005)	3.968
51	Adam(learning_rate=0.001)	5.61
52	Adam(learning_rate=0.007)	5.272
53	Adam(learning_rate=0.01)	5.396
54	Adam(learning_rate=0.015)	5.281
55	Adam(learning_rate=0.02)	5.48


- Learning rate ranging from **0.0005 to 0.02** were tried.
- **Learnings**
- Reducing the learning rate reduced sharpe ratio, where **more** number of epochs might be necessary
- Increasing the learning rate increased the sharpe ratio till 0.02 cutoff post which, it started decreasing

Ensemble model




Exploring ensemble model								
Model No.	Activation function	Batch size	Number of epochs	Optimizer	Number of layers	Network architecture	Patience	Sharpe Ratio
56	relu in both network	128	30	Adam(learning_rate=0.001)	8 and 6	64,64,32,16,8,4,2,1 and 64,32,16,8,4,1	10	5.35

- Though ensemble model provided higher sharpe ratio, creating an **ensemble model of two DNN made sense when the validation dataset is unknown to get a stable model.**




Time to increase number of layers, neurons, Epoch and Patience




Exploring patience, number of epochs and increasing the number of layers and neurons									
Model No.	Activation function	Batch size	Number of epochs	Optimizer	Number of layers	Network architecture	Patience	Sharpe Ratio	
57	relu	128	60	Adam(learning_rate=0.001)	7	256,128,64,32,16,8,4,2,1	15	5.64	
58	relu	128	70	Adam(learning_rate=0.001)	8	512,256,128,64,32,16,8,4,2,1	15	4.75	


- Given other parameters fixed, increasing the number of layers, neurons and subsequently higher epoch and patience to train better increased the sharpe ratio significantly to **5.64**



Final Model – Sharpe ratio 5.64




Best vs Worst Model comparison



Exploring patience, number of epochs and increasing the number of layers and neurons								
Model No.	Activation function	Batch size	Number of epochs	Optimizer	Number of layers	Network architecture	Patience	Sharpe Ratio
57	relu	128	60	Adam(learning_rate=0.001)	7	256,128,64,32,16,8,4,2,1	15	5.64
58	relu	128	70	Adam(learning_rate=0.001)	8	512,256,128,64,32,16,8,4,2,1	15	4.75


S.No	Activation function	Batch size	Number of epochs	Optimizer	Number of layers	Network architecture	Patience	Sharpe Ratio
23	relu	128	30	SGD(learning_rate=0.001, momentum=0.0, nesterov=True)	8	64,64,32,16,8,4,2,1	10	0.77

- Number of layers and number of neurons in each layer, Optimizer, number of epochs play a significant role in model performance




Out of box trials and takeaways

- Given the regression problem, **Activation function was removed in last but one layer**, but did not impact significantly in performance
- Ensemble of two Neural Networks** might be more suitable to produce a robust model when validation set is unknown.
- Alternate and random removal of batch normalization** did not impact the model performance significantly
- Alternate activation function** were tried but did not produce significant improvement in performance when compared to relu



Overall takeaways

- Number of layers and neurons are the major parameter** which produces significant improvement in the model performance
- Swish activation function performed significantly better apart from relu.** More epochs might improve the prediction
- After a threshold, it is **not necessary that if we increase the number of layers and epochs to very high value it will improve accuracy**



Dataset 2



Number of layers and Activation Functions

Number of Layers			
Model no.	Number of layers	Number of neurons	R2
1	6	2056,256,128,64,32,1	-1.72
2	7	2056,1024,512,256,128,64,1	-1.29
3	8	2056,1024,512,256,128,64,32,1	-0.534073
4	9	4098,2056,1024,512,256,128,64,32,1	-1.34

- **Learning** – With increase in layers, the R2 got better only till 8 layers. It started overfitting after that

Activation Function			
Model no.	Number of layers	Activation functions	R2
5	8	tanh	-0.56
6	8	Relu	-1.93
7	8	Prelu	-1.27
8	8	sigmoid	-0.37
9	8	Sigmoid and tanh	-0.302594

- Individually, Sigmoid and tanh performed very well. One interesting finding is adding sigmoid and tanh alternately improved the R2

Number of Dropout layers and Loss Function

- Since it was a regression problem what if we add two neuron layer before the last layer!

Dropout ratio and Number of Dropout			
Model no.	Number of layers	Number of neurons	R2
10	3	0.2,0.2,0.2	-0.40
11	4	0.3,0.3,0.3,0.05	-0.312615
12	5	0.3,0.3,0.30,0.1,0.05	-0.40

- Learning** - It is very evident about the importance of dropout layers as this could reduce the overfitting the data

Loss Function			
Model no.	Loss function	R2	
13	Mean absolute error	-20.858745	
14	an squared logarithmic e	-0.454134	
15	mse	-0.250568	


- Learning** - Solving the loss function with mse as hyperparameters gives the best model

Number of Epochs, Batch size

Epochs			
Model no.	Epochs	Batch size	Sharpe Ratio
16	100	30	-35.22
17	100	30	-0.32
18	120	20	-0.46
19	120	20	-0.33
20	120	40	-0.300459


- **Learning** – With increase in Epochs, the R2 is improving but it is very computationally heavy

Optimizer





Model no.	Optimizer	Sharpe Ratio
20	SGD	-20.86
21	Adam	-0.45
22	RMSprop(learning_rate=0.005, momentum=0.9, nesterov=False)	-0.250568
23	RMSprop(learning_rate=0.001, momentum=0.9, nesterov=False)	-0.28
24	RMSprop(learning_rate=0.004, momentum=0.9, nesterov=True)	-0.63


- **Learning** – RMSprop with the momentum of 0.9 is the best among all the optimizers with the existing architecture




Best Model



Optimizer = RMS Prop(learning_rate=0.0005, momentum=0.9),
loss='mean_absolute_error', patience=30, batch_size=32, epochs=100



Exploding Gradient Descent




Best Model vs Least Accurate

		Best Model	Least Accurate Model
S.No.		22	14
R2		-0.250568	-35.217213
Optimizer	Optimizer	RMSprop	adam
	Rate	0.0005	Default
	Momentum	0.9	Default
Loss		mean_absolute_error	mse
Epochs		100	100
Batch size		32	32
Patience		30	20
verbose		1	1
Layer1	Neurons	Layer Normalization	Layer Normalization
	Activation Function		
Layer2	Neurons	4098	4098
	Activation Function	tanh	relu
Layer3	Neurons	Batch Normalization	Batch Normalization
	Activation Function		
Layer4	Neurons	DropOut 0.5	DropOut 0.5
	Activation Function		
Layer5	Neurons	2056	2056
	Activation Function	sigmoid	relu
Layer6	Neurons	Batch Normalization	Batch Normalization
	Activation Function		
Layer7	Neurons	1024	1024
	Activation Function	tanh	relu
Layer8	Neurons	DropOut 0.2	DropOut 0.2
	Activation Function		

Layer9	Neurons	Batch Normalization	Batch Normalization
	Activation Function		
Layer10	Neurons	512	512
	Activation Function	sigmoid	relu
Layer11	Neurons	Batch Normalization	Batch Normalization
	Activation Function		
Layer12	Neurons	DropOut 0.1	DropOut 0.1
	Activation Function		
Layer13	Neurons	256	256
	Activation Function	tanh	relu
Layer14	Neurons	BatchNormalization	BatchNormalization
	Activation Function		
Layer15	Neurons	DropOut 0.1	DropOut 0.1
	Activation Function		
Layer16	Neurons	128	128
	Activation Function	sigmoid	relu
Layer17	Neurons	BatchNormalization	BatchNormalization
	Activation Function		
Layer18	Neurons	Dropout 0.05	Dropout 0.05
	Activation Function		
Layer19	Neurons	64	64
	Activation Function	tanh	relu
Output Layer		Output1	Output1

Learnings



Overall Learnings

1. It is very important to use our intuition while training the deep neural networks as it is not practical to use all the possible combinations


2. There are new combinations that have tried:
 - a. Kernel Initializer
 - b. Encountered Vanishing Gradient descent
 - c. Encountered Exploding Gradient descent
 - d. Importance of Dropout
 - e. Alternate activation functions

3. The Modelling exercise is all about trying different parameters and understanding the directions

"We learned a lot about what not to do!!"

- Deep Learning thoughts

Tweaking Neural Net

A cartoon illustration of Peter Griffin, the father from the TV show Family Guy. He is a middle-aged man with a large belly, wearing a white button-down shirt and green pants. He is standing in front of a window with horizontal blinds, looking out with a thoughtful expression while holding a pencil. The background shows a dark room and a window frame.

Parameters